

TOPIC FOUR

COMPUTER VISION

SUMMARY NOTES (ET0732)

Lecture Notes: Introduction to Computer Vision Fundamentals with Google Cloud

1. Overview of Computer Vision and Its Rapid Growth

- **Definition:** Computer Vision is a subset of Machine Learning (ML) and Artificial Intelligence (AI) that enables computers to interpret and understand visual data from images and videos.
 - **Historical Context:**
 - Origin of digital imaging in the 1960s.
 - Significant advancements over the last few decades, driven by cloud computing and specialized hardware.
 - **Growth Factors:**
 - Increasing high-resolution image data, thanks to the proliferation of smartphones.
 - Massive amounts of visual data generated, e.g., 79 zettabytes of data in 2021.
 - Internet usage: 93% of people accessed the web through mobile devices.
 - **Scale of Data:**
 - Every second: 8.5 hours of videos uploaded to YouTube.
 - Google Photos stored 4 trillion photos by November 2020, with 28 billion new photos and videos uploaded weekly.
-

2. Understanding Computer Vision Problems

- **Types of Problems:**
 - **Image Classification:** Assigns labels to whole images (e.g., identifying an object category).
 - **Semantic Segmentation:** Labels each pixel, segmenting an image into different categories.
 - **Instance Segmentation:** Differentiates objects in an image by identifying boundaries and coloring pixels uniquely.
 - **Image Classification with Localization:** Labels an image and provides bounding boxes for object locations.

- **Object Recognition:** Identifies objects and provides class labels with probabilities but not locations.
 - **Object Detection:** Detects and locates multiple objects, providing bounding boxes.
 - **Pattern Recognition:** Identifies repeated shapes, colors, and patterns (e.g., facial recognition, OCR).
 - **Edge Detection:** Finds object boundaries by detecting changes in brightness and intensity.
 - **Feature Matching:** Compares features across images, enabling applications like object tracking and 3D reconstruction.
-

3. Historical Developments in Computer Vision

- **Early Research:**
 - 1960s: Studies on cat vision laid the foundation.
 - 1974: First robust Optical Character Recognition (OCR) system developed, shifting focus to practical applications.
 - **2000s Onwards:** Advances in handling complex tasks like image segmentation, object detection, and facial recognition.
-

4. Business Applications of Computer Vision

- **Examples:**
 - **The New York Times:** Uses machine learning for digitizing and categorizing vast photo archives.
 - **Box:** Employs image labeling technology for efficient searching and content analysis.
 - **Harte Research Institute at Texas A&M University:** Uses computer vision to classify shorelines from aerial imagery for environmental analysis.
-

5. Google Cloud's Machine Learning Tools

- **Pre-Built Vision API:**
 - Analyzes images to detect objects, text, emotions, and more.
 - Outputs ranked labels, bounding boxes, and even generates descriptive captions.
- **Capabilities:**
 - Classifies and segments images, detects facial attributes, recognizes text, and automates tasks.

- Example: Describes images (e.g., “Two hockey players are fighting over a puck”) and may occasionally make errors.
-

6. Potential and Limitations of Computer Vision Models

- **Applications:**
 - Object detection, disease diagnosis, autonomous driving, etc.
 - **Performance:** Some models outperform humans in specific tasks.
 - **Challenges:** Even advanced models can misinterpret images, highlighting the need for continuous improvement.
-

Lecture Notes: Custom Training with Linear, Neural Network, and Deep Neural Network Models

Overview

- **Goal:** Classify images using machine learning techniques, exploring their limitations with image datasets.
 - **Focus:** Create custom image classification models from scratch without pre-trained weights, starting with the 5-flowers dataset.
-

Image Dataset: 5-Flowers

- **Dataset:** 3,700 labeled photographs of five types of flowers.
- **Image Classification Problem:** Classify images into one of the five categories (multiclass classification).

Image Representation

- **Images:** Represented as pixels (2D arrays of numbers) and fed into models as inputs.
- **Classes:** Five possible output classes for classification.

Training and Evaluation

- **Data Split:** Training set and test set.
 - **Storage:** Data is stored in Cloud Storage and accessed using TensorFlow datasets.
-

Techniques for Model Training

1. **Linear Models:** Assumes a linear relationship between input and output.
 - Uses softmax/sigmoid for classification.
 - Computes a weighted sum of inputs plus a bias term.

2. Neural Network Models:

- Extends linear models with multiple layers.
- Incorporates regularization techniques like dropout to prevent overfitting.
- Batch normalization to stabilize and speed up training.

3. Deep Neural Networks:

- Multiple hidden layers for learning complex representations.
- Dropout and batch normalization enhance model performance.

Concepts

- **Dropout:** Reduces overfitting by randomly setting activations to zero during training.
 - **Batch Normalization:** Normalizes inputs to layers, improving stability and training speed.
-

Image Data Processing Workflow

1. Input Pipeline Setup:

- Use `tf.io` and `tf.image` for processing images.
- Read, decode, convert, and resize images to prepare them for the model.

2. Data Conversion:

- Images are read as byte tensors and decoded into 3D uint8 tensors.
- Convert RGB values to floats and scale between 0 and 1.

3. Data Resizing:

- Resize images as needed using `tf.image.resize`.
- Consider padding or cropping to maintain aspect ratio.

4. Visualization:

- Use Matplotlib to visualize images and understand data distribution.
-

Efficient Data Handling

- **Full Dataset Reading:** Use `tf.data.Dataset` API to create efficient input pipelines.
 - **Transformations:** Preprocess data in a streaming fashion for memory efficiency.
 - **Batching:** Batch data for model optimization and efficient computation.
-

Data Extraction and Labeling

- **Class Names:** Extracted from image filenames or provided CSV files.

- **Data Parsing:** Use TensorFlow functions to streamline data preparation and label extraction.

Creating the Dataset

- **Methods:** Use TextLineDataset or TFRecordDataset for structured data processing.
- **Function Definition:** Parse CSV lines, extract filenames, and read images with labels.

Lecture Notes on Implementing Linear Models and Deep Neural Networks with Keras

Introduction to Linear Models

- **Start Simple:** Implement simpler models first, introducing complexity only if necessary to meet performance criteria.
 - Simple models are less prone to overfitting and are easier to interpret and maintain.
- **Alternative Approach:** Some techniques use large models with strong regularization from the start.

Implementing Linear Models Using Keras

- **Keras:** High-level API of TensorFlow, designed for efficient and intuitive model building, focusing on modern deep learning.
- **Core Concepts:**
 - **Layers and Models:** Fundamental structures in Keras.
 - **Sequential Model:** A straightforward stack of layers where each layer feeds into the next. Suitable for simple, linear models but not for models needing multiple inputs/outputs or complex topologies.

Steps to Create and Train a Model in Keras

1. **Define and Create Model:** Use the Sequential class and stack layers.
2. **Compile Model:** Configure optimizer, loss function, and metrics.
3. **Optional:** Use model.summary() to display the model architecture.
4. **Train Model:** Use model.fit() with training data, specifying the number of epochs.
5. **Evaluate Model:** Plot training and validation loss/accuracy to assess performance.
6. **Make Predictions:** Use model.predict() and convert logits to probabilities using softmax.

Example Model Architecture

```
python
```

Copy code

```
model = tf.keras.Sequential([  
    tf.keras.layers.Flatten(input_shape=(IMG_HEIGHT, IMG_WIDTH, 3)),  
    tf.keras.layers.Dense(len(CLASS_NAMES))  
])
```

- **Flatten Layer:** Converts a 3D image (e.g., 224x224x3) into a 1D array. It doesn't learn parameters; it only reshapes data.
 - **Dense Layer:** Computes the weighted sum of inputs and applies an activation function like softmax to output probabilities.
-

Model Compilation

- **Settings:**
 - **Optimizer:** E.g., Adam, adjusts model weights to minimize loss.
 - **Loss Function:** Measures the model's error; SparseCategoricalCrossentropy is used for multi-class classification.
 - **Metrics:** E.g., accuracy, to monitor performance.
-

Model Training and Evaluation

- **Training:** Use model.fit() with both training and validation datasets.
 - **History and Callbacks:** Record and plot loss and accuracy to analyze the model's convergence and detect overfitting.
-

Making Predictions and Understanding Outputs

- **Model Outputs:** Logits transformed into probabilities via softmax.
 - **Confidence Values:** Higher confidence indicates stronger model prediction for a class.
-

Deep Neural Networks (DNNs) and Regularization Techniques

- **Universal Approximation Theorem:** A single hidden layer can theoretically solve any problem, but large models are impractical.
- **Challenges with Large Models:**
 - Increased memory usage, slower training, and higher risk of overfitting.
- **Overfitting Prevention:**

- **Regularization:** Dropout, L1, and L2.
 - **Dropout:** Randomly deactivates neurons during training to prevent overfitting.
 - **Batch Normalization:** Normalizes inputs of each layer to speed up training and improve stability.
-

Dropout and Batch Normalization

- **Dropout:** Applies a dropout layer with a probability p to deactivate neurons during training.
 - **Batch Normalization:** Normalizes activations using batch statistics, stabilizing and accelerating learning.
-

Implementing Regularization in Keras

- **Dropout:** `tf.keras.layers.Dropout(rate)`, used only during training.
 - **Batch Normalization:** `tf.keras.layers.BatchNormalization()`, applied before the activation function.
-

Summary

- **Model Building:** Learned how to represent and train linear models, and progressively more complex neural networks for image classification.
 - **Regularization Techniques:** Dropout and batch normalization to improve model generalization.
-

Lecture Notes: Introduction to Convolutional Neural Networks (CNNs)

1. Overview of CNNs

- CNNs specialize in detecting visual patterns within images.
 - They differ from traditional neural networks by employing **convolutions** to extract features, sliding filters over image pixels to imitate how the human brain processes visual data.
-

2. Key Differences and Concepts

- **Feature Extraction:** Convolutions mimic the hierarchical approach of the visual cortex, emphasizing locally correlated features.

- CNNs use **filters** (kernels) that operate on images to create **feature maps**.
 - **Parameters to Consider:**
 - **Filters:** The number and type affect feature extraction.
 - **Channels:** Define input data (e.g., RGB for color images).
 - **Kernel Size:** Determines the receptive field of filters.
 - **Strides & Padding:** Control how filters slide and manage border effects.
 - **Activation Functions:** Introduce non-linearity, such as ReLU.
-

3. Pooling Layers

- **Purpose:** Reduce the spatial dimensions of feature maps, making computations more efficient and reducing sensitivity to object location.
 - Common types include **Max Pooling** (selects maximum value) and **Average Pooling** (computes average value).
-

4. Evolution and History

- **1980:** Kunihiro Fukushima developed the **Neocognitron**, inspired by simple and complex cells in the human brain.
 - **Late 1980s-1990s:** CNNs were formalized by **Yann LeCun**, leading to models like **LeNet**, which was pivotal for tasks like handwriting recognition.
 - **2012:** **AlexNet**'s win in the ImageNet competition marked a breakthrough, establishing CNNs as a key framework in computer vision.
-

5. CNN Applications and Beyond

- Widely used for **image classification** and **object detection**.
 - Can also be applied to non-image data (e.g., audio, time series) but rely heavily on local feature hierarchies.
 - Prior to CNNs, feature engineering required manual image preprocessing, whereas CNNs learn relevant features autonomously.
-

6. Dense Layers Recap

- **Dense Layers:** Every input is connected to every neuron, leading to vast parameter counts for high-resolution images.
- Dense layers lack hierarchical structuring, making pixel order irrelevant for classification tasks.

- CNNs contrast this by preserving the local structure, vital for visual tasks.
-

7. Feature Engineering and Modern CNNs

- Traditional models relied on engineered features (e.g., using Gabor filters) for pattern recognition.
 - **Post-2012:** CNNs simplify design by learning features directly from images, with models like **AlexNet** and **Inception** setting standards.
-

8. Detailed CNN Layers

- **Convolutional Layers:** Use kernels to detect spatial patterns, with shared weights promoting efficiency.
 - **Example:** A 5x5 image processed with a 3x3 kernel results in a smaller feature map, reduced by kernel size minus one.
 - CNNs stack multiple filters, each detecting different features.
-

9. Specialized CNN Types

- **1D CNNs:** For sequential data (e.g., audio, time series).
 - **3D CNNs:** For 3D data (e.g., videos, MRI scans).
 - **Input Representation:** Grayscale images use 2D tensors; color images use 3D tensors (RGB channels).
-

10. Kernel Operations

- Kernels detect edges and patterns by computing the weighted sum of pixels.
 - Example: Edge detection kernels highlight areas with intensity changes.
 - **Parameter Sharing:** Kernels have consistent weights, enhancing efficiency.
-

11. Implementing CNNs

- **Keras** simplifies CNN construction, automating layer setup and enabling efficient feature learning.
 - **Edge Detection Example:** Using weighted kernels to identify horizontal and vertical features.
 - CNNs learn filters to progressively build complex representations, from edges to entire objects.
-

12. Practice and Quiz Insights

- **Quiz Concept:** Understanding kernel operations and how convolutions identify features.
- Practical implementation involves configuring filters, activation functions, and other parameters.

Lecture Notes: CNN Model Parameters and Operations

1. ML Model Parameters Recap

- **Definition:** Parameters are values learned during training to transform input data into the predicted output.
 - **Complexity and Parameters:** As a model's complexity increases, so does the number of trainable parameters.
 - *Example:* Simple neural networks have weights and biases; CNNs have additional parameters.
-

2. Creating a Convolution Layer in Keras

- **Conv2D Layer:** `tf.keras.layers.Conv2D` method creates a 2D convolutional layer in Keras.
 - **Input/Output Format:** CNNs handling image recognition expect 4D tensors: [batch, height, width, channels].
 - *Example:* For 256x256 RGB images, the input shape is [256, 256, 3]. With a batch of 16, it becomes [16, 256, 256, 3].
-

3. CNN Model Parameters

- **Number of Filters:** Determines the number of independent filters applied. Output channels = number of filters.
- **Input Channels:** Based on the input image's channel count. For a 256x256 RGB image: 3 channels.
- **Kernel Size:** Defines the dimensions of each filter (e.g., 3x3, 5x5). Smaller kernels with multiple layers are efficient.
- **Strides:** Step size for the filter sliding across the image. Default is 1. Larger strides reduce output size but can skip information.
- **Padding:** Adds borders to maintain input-output dimensions. Methods:
 - `same`: Keeps output size the same as input.

- valid: No padding; reduces output size.
-

4. Calculating Parameters in CNNs

- **Convolution Layer:** $\text{Parameters} = (\text{width} * \text{height} * \text{input channels} + 1) * \text{number of filters}$.
 - **Pooling Layers:** No learnable parameters; used to reduce dimensionality and computation.
 - **Fully Connected Layers:** High parameter count. $\text{Parameters} = (\text{neurons in current layer} * \text{neurons in previous layer} + 1) * \text{current neurons}$.
-

5. Pooling Operations

- **Purpose:** Reduce feature map dimensions and computations.
 - **Max Pooling:** Returns the maximum value within a filter window.
 - *Example:* 4x4 input reduced to 2x2 using a 2x2 filter with a stride of 2.
 - **Average Pooling:** Computes average values in the filter window.
 - **Global Pooling:** Used for summarizing information from the feature map.
 - **Implementation in Keras:** `tf.keras.layers.MaxPooling2D(pool_size=2, strides=1)`
-

6. Convolution vs. Dense Layers

- **Dense Layers:** Connect all input pixels to every neuron, requiring a large number of parameters.
 - **Convolution Layers:** Use kernels to detect patterns (edges, textures) with fewer weights, improving efficiency and reducing training time.
 - **Example:** MNIST dataset - using convolution reduces parameter count compared to dense layers.
-

7. Key Takeaways

- CNNs use fewer parameters than dense layers, enhancing efficiency in image processing tasks.
- Convolution and pooling layers recognize and compress patterns, passing a flattened feature vector to a fully connected network.
- Use `model.summary()` in Keras to check parameter details easily.

Quiz Question Review

- **Advantage of Convolution Kernels:** Reduces parameters by processing small patches and sharing weights, speeding up training.

Lecture Notes: Dealing with Image Data in Computer Vision

1. Overview of Section:

- **Objective:** Learn how to preprocess image data for reproducibility in production.
 - **Hands-On Labs:** Implement preprocessing in Keras using TensorFlow datasets.
 - **Key Topics:** Relationship between model parameters and data scarcity, data augmentation, transfer learning, and efficient data storage formats.
-

2. Creating and Reading Image Datasets:

- **Creating a Dataset:** Collect and annotate image data using internal resources or third-party services.
 - **Validating Data:** Ensure image data is in formats like JPEG/PNG but avoid inefficiencies (e.g., CSV files) for large datasets.
 - **Efficient Data Format:** Use tf.records for better I/O operations and storage efficiency, essential for optimized training on frameworks like TPUs.
-

3. TensorFlow Data Handling:

- **Using tf.data API:** Construct efficient data pipelines with TFRecordDataset for fast processing.
 - **Example Pipeline:** Aggregate data, apply random transformations, and batch images for training.
-

4. Image Preprocessing:

- **Purpose:** Prepare raw images for model training by resizing, color conversion, and other transformations.
 - **Techniques:** Resizing (using tf.image.resize), flipping, rotating, and cropping images to match model input shapes.
 - **Aspect Ratio Management:** Use tf.image.resize_with_pad to maintain the original aspect ratio and avoid distortion.
 - **Learned Resizers:** Enhance performance using custom resizers beyond traditional methods like bilinear.
-

5. Keras Preprocessing Layers:

- **API Usage:** Keras provides layers like Resizing and Rescaling to standardize inputs.

- **Integration:** Combine preprocessing layers with models to streamline training and prediction.
-

6. Handling Data Scarcity:

- **Problem Overview:** Modern models require vast labeled datasets, but data can be limited.
 - **CNNs' Needs:** Models with more parameters demand more data for effective training.
-

7. Strategies for Data Scarcity:

- **Data Augmentation:** Increase dataset size by applying transformations (e.g., cropping, flipping, color adjustments).
 - **Transfer Learning:** Use pre-trained models to leverage knowledge from similar tasks, reducing data requirements.
-

8. Data Augmentation Techniques:

- **Common Transformations:** Blurring, sharpening, resizing, rotating, flipping, and color adjustments.
 - **Considerations:** Ensure transformations do not compromise the model's learning. E.g., orientation is critical in cases like distinguishing flags or species.
 - **Task-Specific Augmentation:** Tailor techniques to the data's characteristics and domain requirements.
-

9. Implementing Data Augmentation:

- **Using `tf.image`:** Write augmentation pipelines with methods like `flip_left_right` or `rgb_to_grayscale`.
 - **Random Transformations:** Use APIs like `tf.image.stateless_random_brightness` for controlled randomness.
 - **Parallelization with `tf.data`:** Use `Dataset.map` to parallelize preprocessing.
-

10. Keras Image Augmentation Layers:

- **Built-In Layers:** `RandomTranslation`, `RandomRotation`, `RandomZoom`, etc., for augmenting images during training.
- **Inference Handling:** Preprocessing is only applied during training, simplifying the prediction phase.
- **Training Considerations:** Training with augmented data often requires longer training times due to increased data size.

11. Quiz Questions (Examples):

- **Parameter Count Calculations:** Analyze the parameter count in models, e.g., linear and convolutional layers.
- **Practical Applications:** Discuss scenarios where specific augmentations would or would not improve model performance.

Lecture Notes on Transfer Learning

1. Introduction to Transfer Learning

- **Purpose:** A method to address data scarcity by decreasing the need for a large amount of labeled data.
 - **Approach:** Instead of creating more data, transfer learning initializes model parameters with values from a pre-trained model, enhancing efficiency.
-

2. Optimization and Training

- **Optimization Journey:** The process of finding optimal weights for the model to minimize data and time expense.
 - **From Scratch vs. Transfer Learning:**
 - *Training from Scratch:* Time-consuming and resource-intensive.
 - *Transfer Learning:* Utilizes a pre-trained model trained on similar tasks, saving both time and data.
-

3. Core Concept of Transfer Learning

- **How It Works:** Knowledge from a source model trained on a large dataset is transferred to a new, related task.
 - **Benefit:** Significantly reduces training time compared to starting from scratch, especially when source and target tasks are similar.
-

4. Example: Using ImageNet for Transfer Learning

- **ImageNet:** A large dataset with 14 million labeled images across 20,000 categories.
- **Naive Transfer Learning:** Using the ImageNet model directly for predictions may not work if the target classes differ in number or specificity.
- **Advanced Transfer Learning:** Involves modifying the source model:
 - Replace parts closely tied to the source task (e.g., output layers).

- Retain generalized parts (e.g., convolutional layers) for feature extraction.
-

5. Model Layers and Task Dependence

- **Input Layer:** Task-independent and can handle general input, like any RGB image.
 - **Convolutional Layers:** Generally task-independent, used for feature extraction.
 - **Output Layers:** Highly task-dependent, aligned with the specific output requirements of the source task.
 - **Feature Hierarchy:**
 - CNNs learn general to specific features.
 - Early layers detect simple patterns; later layers become more specific and task-dependent.
-

6. Adjusting the Source Model

- **Where to Cut the Network:** No clear point due to distributed representations in neural networks.
 - **Standard Practice:** Cut after the convolutional layers and add fully connected layers suited to your task.
 - **Weight Training Decisions:**
 - *Constant Weights:* Use the source model as a feature extractor, recommended for small datasets to avoid overfitting.
 - *Trainable Weights:* Adjust weights if your dataset is large enough, reducing overfitting risks.
-

7. Implementing Transfer Learning

- **Pre-trained Models:** Models trained on extensive datasets (e.g., ImageNet) that are readily available for use.
 - **MobileNet Example:**
 - Pre-trained on ImageNet with 1-4 million parameters.
 - Efficient at compressing visual information, making it suitable for related image classification tasks.
 - Available on TensorFlow Hub for easy integration with Keras.
-

8. Recap and Applications

- **Preprocessing for Reproducibility:** Important for consistent performance in production.

- **Data Augmentation vs. Transfer Learning:**
 - *Data Augmentation*: Increases data variety and model robustness.
 - *Transfer Learning*: Leverages existing knowledge to minimize data needs.
- **Next Steps**: Reviewing key concepts and applying them to practical problems.